

# FastVPINNs: An efficient tensor-based Python library for solving partial differential equations using hp-Variational Physics Informed Neural Networks

Thivin Anandh<sup>1¶</sup>, Divij Ghose<sup>1</sup>, and Sashikumaar Ganesan<sup>1¶</sup>

<sup>1</sup> Department of Computational and Data Sciences, Indian Institute of Science, Bangalore, India ¶ Corresponding author

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

## Software

- Review [↗](#)
- Repository [↗](#)
- Archive [↗](#)

Editor: [Open Journals](#) [↗](#)

Reviewers:

- [@openjournals](#)

Submitted: 01 January 1970

Published: unpublished

## License

Authors of papers retain copyright and release the work under a

Creative Commons Attribution 4.0 International License ([CC BY 4.0](#))

## Introduction

Partial differential equations (PDEs) are essential in modeling physical phenomena such as heat transfer, electromagnetics and fluid dynamics. The current progress in the field of scientific machine learning (SciML) has resulted in incorporating deep learning and data-driven methods to solve PDEs. Let us consider a two-dimensional Poisson equation as an example, defined on the domain  $\Omega$ , which reads as follows:

$$-\nabla^2 u(x) = f(x), \quad \Omega \in \mathbb{R}^2, \quad (1)$$

$$u(x) = g(x), \quad x \in \partial\Omega. \quad (2)$$

Here,  $x \in \Omega$  is the spatial co-ordinate,  $u(x)$  is the solution of the PDE,  $f(x)$  is a known source term and  $g(x)$  is the value of the solution on the domain boundary,  $\partial\Omega$ .

A neural network is a parametric function of  $x$ , denoted as  $u_{\text{NN}}(x; W, b)$ . In this context,  $W$  and  $b$  represent the weights and biases of the network. When the neural network consists of  $h$  hidden layers, with the  $i^{\text{th}}$  layer containing  $n_i$  neurons, the mathematical representation of the function takes the following shape:

$$u_{\text{NN}}(x; W, b) = l \circ T^{(h)} \circ T^{(h-1)} \dots \circ T^{(1)}(x)$$

Here,  $l: \mathbb{R}^{n_h} \rightarrow \mathbb{R}$  is a linear mapping in the output layer and  $T^{(i)}(\cdot) = \sigma(W^{(i)} \cdot + b^{(i)})$  is a non-linear mapping in the  $i^{\text{th}}$  layer ( $i = 1, 2, \dots, h$ ), with the non-linear activation function  $\sigma$  and the weights  $W^{(i)}$  and biases  $b^{(i)}$  of the respective layers.

Physics-informed neural networks (PINNs), introduced by (Raissi et al., 2019), work by incorporating the governing equations of the physical systems and the boundary conditions as physics-based constraints to train the neural networks. The core idea of PINNs lies in the ability of obtaining the gradients of the solutions with respect to the input using the automatic differentiation routines available within deep learning frameworks such as tensorflow (Abadi et al., 2015). The loss function for the PINNs can be written as follows:

$$L_p(W, b) = \frac{1}{N_T} \sum_{t=1}^{N_T} |(-\nabla^2 u_{\text{NN}}(x_t; W, b) - f(x_t))|^2,$$

$$L_b(W, b) = \frac{1}{N_D} \sum_{d=1}^{N_D} |u_{\text{NN}}(x_d; W, b) - g(x_d)|^2,$$

$$L_{\text{PINN}}(W, b) = L_p + \tau L_b.$$

Here, the output of the neural network,  $u_{\text{NN}}(x; W, b)$ , is used to approximate the unknown solution. In addition,  $N_T$  is the total number of training points in the interior of the domain  $\Omega$ ,  $N_D$  is the total number of training points on the boundary  $\partial\Omega$ ,  $\tau$  is a scaling factor applied to control the penalty on the boundary term and  $L_{\text{PINN}}(W, b)$  is the loss function of the PINNs.

Variational Physics informed neural networks (VPINNs) (Kharazmi et al., 2019) are an extension of PINNs, where the weak form of the PDE is used in the loss function of the neural network. The weak form of the PDE is obtained by multiplying the PDE with a test function, integrating over the domain and applying integration by parts to the higher order derivative terms. The method of hp-VPINNs (Kharazmi et al., 2021) was subsequently developed to increase the accuracy using h-refinement (increasing number of elements) and p-refinement (increasing the number of test functions). The domain  $\Omega$  is divided into an array of non-overlapping cells, labeled as  $K_k$ , where  $k = 1, 2, \dots, N_{\text{elem}}$ , ensuring that the complete union  $\bigcup_{k=1}^{N_{\text{elem}}} K_k = \Omega$  covers the entire domain  $\Omega$ . The hp-VPINNs framework utilizes specific test functions  $v_k$ , where  $k$  ranges from 1 to  $N_{\text{elem}}$ , that are localized and defined within individual non-overlapping element across the domain.

$$v_k = \begin{cases} v^p \neq 0, & \text{over } K_k, \\ 0, & \text{elsewhere.} \end{cases}$$

The loss function of hp-VPINNs with  $N_{\text{elem}}$  elements in the domain can be written as follows.

$$L_v(W, b) = \frac{1}{N_{\text{elem}}} \sum_{k=1}^{N_{\text{elem}}} \left| \int_{K_k} \left( \nabla u_{\text{NN}}(x; W, b) \cdot \nabla v_k dK - \int_{K_k} f v_k dK \right) \right|^2,$$

$$L_b(W, b) = \frac{1}{N_D} \sum_{d=1}^{N_D} |u_{\text{NN}}(x; W, b) - g(x)|^2,$$

$$L_{\text{VPINN}}(W, b) = L_v + \tau L_b.$$

Where  $K_k$  is the  $k^{\text{th}}$  element in the domain,  $v_k$  is the test function in the respective element. Further,  $L_v(W, b)$  is the weak form PDE residual and  $L_{\text{VPINN}}(W, b)$  is the loss function of the hp-VPINNs. For more information on the derivation of the weak form of the PDE and the loss function of hp-VPINNs, refer to (Anandh et al., 2024) and (Ganesan & Tobiska, 2017)

## Statement of need

The existing implementation of hp-VPINNs framework (Kharazmi, 2023) suffers from two major challenges. One being the inability of the framework to handle complex geometries and the other being the increased training time associated with the increase in number of elements within the domain. In the work (Anandh et al., 2024), we presented FastVPINNs, which addresses both of these challenges. FastVPINNs handles complex geometries by using bilinear transformation, and it uses a tensor-based loss computation to reduce the dependency of training time on number of elements. The current implementation of FastVPINNs can achieve a speed-up of upto a 100 times when compared with the existing implementation of hp-VPINNs. We have also shown that with proper hyperparameter selection, FastVPINNs can outperform PINNs both in terms of accuracy and training time, especially for problems with high frequency solutions.

Our FastVPINNs framework is built using TensorFlow-v2.0 (Abadi et al., 2015), and provides an elegant API for users to solve both forward and inverse problems for PDEs like the Poisson, Helmholtz and Convection-Diffusion equations. With the current level of API abstraction, users should be able to solve PDEs with less than six API calls as shown in the minimal working example section. The framework is well-documented with examples, which can enable users to get started with the framework with ease.

66 The ability of the framework to allow users to train a hp-VPINNs to solve a PDE both faster  
67 and with minimal code, can result in widespread application of this method on several real-world  
68 problems, which often require complex geometries with a large number of elements within the  
69 domain.

70 **Modules**

71 The FastVPINNs framework consists of five core modules, which are Geometry, FE, Data,  
72 Physics and Model.

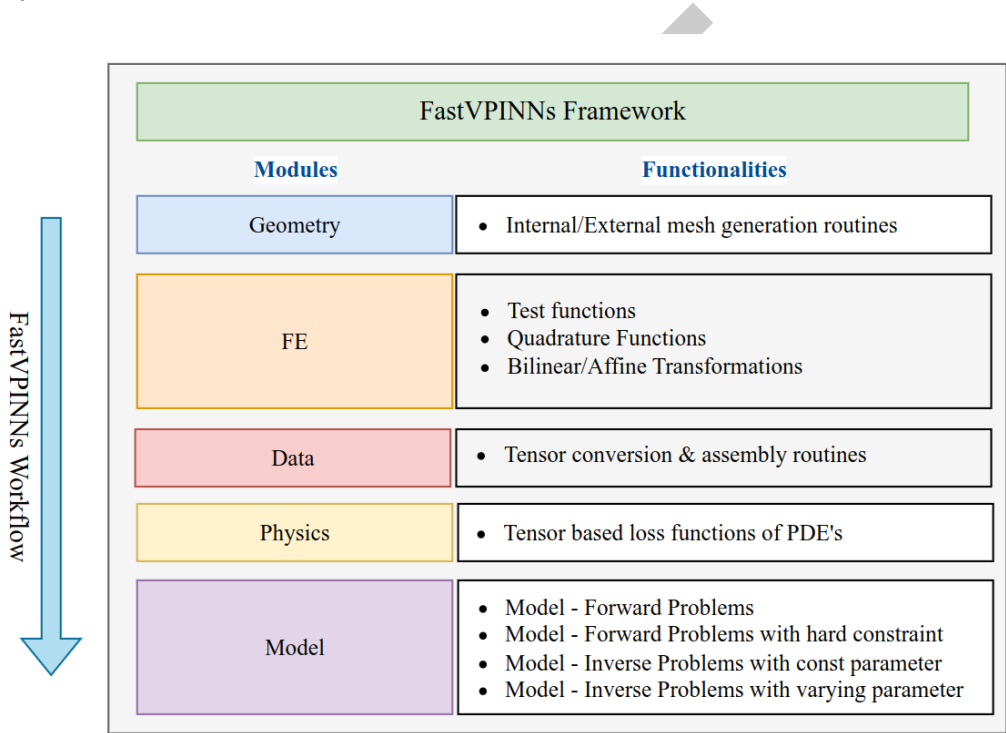


Figure 1: FastVPINNs Modules

73 **Geometry Module**

74 This module provides the functionality to define the geometry of the domain. The user can  
75 either generate a quadrilateral mesh internally or read an external .mesh file. The module also  
76 provides the functionality to obtain boundary points for complex geometries.

77 **FE Module**

78 The FE module is responsible for handling the finite element test functions and their gradients.  
79 The module's functionality can be broadly classified into into four categories:

- 80     ▪ **Test functions:** Provides the test function values and its gradients for a given reference  
81       element. In our framework, we use Lagrange polynomials as the test functions.
- 82     ▪ **Quadrature functions:** Provides the quadrature points and weights for a given element  
83       based on the quadrature order and the quadrature method selected by the user. This  
84       method will be used for performing the numerical integration of the weak form residual  
85       of the PDE.

- 86     ▪ **Transformation functions:** Provides the implementation of transformation routines such  
87     as affine transformation and bilinear transformation. This can be used to transform the  
88     test function values and gradients from the reference element to the actual element.
- 89     ▪ **Finite Element Setup:** Provides the functionality to set up the test functions, quadrature  
90     rules and the transformation for every element and save them in a common class to  
91     access them. Further, it also hosts functions to plot the mesh with quadrature points,  
92     assign boundary values based on the boundary points obtained from the geometry module  
93     and calculate the forcing term in the residual.

94     *Remark: The module is named FE (Finite Element) Module because of its similarities with*  
95     *classical FEM routines, such as test functions, numerical quadratures and transformations.*  
96     *However, we would like to state that our framework is not an FEM solver, but an hp-VPINNs*  
97     *solver*

## 98     **Data Module:**

99     The Data module collects data from all modules which are required for training and converts  
100     them to a tensor data type with the user specified precision (for example, `tf.float32` or  
101     `tf.float64`). It also assembles the test function values and gradients to form a three-  
102     dimensional tensor, which will be used during the loss computation.

## 103     **Physics Module:**

104     This module contains functions which are used to compute the variational loss function for  
105     different PDEs. These functions accept the test function tensors and the predicted gradients  
106     from the neural network along with the forcing matrix to compute the PDE residuals using  
107     tensor-based operations.

## 108     **Model Module:**

109     This module contains custom subclasses of the `tensorflow.keras.Model` class, which are  
110     used to train the neural network. The module provides the functionality to train the neural  
111     network using the tensor based loss computation and also provides the functionality to predict  
112     the solution of the PDE for a given input.

## 113     **Minimal Working Example**

114     With the higher level of abstraction provided by the FastVPINNs framework, users can solve a  
115     PDE with just six API calls. A Minimal working example to solve the Poisson equation using  
116     the FastVPINNs framework can be found [here](#). The example files with detailed documentation  
117     can be found in the [Tutorials section](#) of the documentation.

## 118     **Testing**

119     The FastVPINNs framework has a strong testing suite, which tests the core functionalities of  
120     the framework. The testing suite is built using the `pytest` framework and is integrated with  
121     the continuous integration pipeline provided by Github Actions. The testing functionalities can  
122     be broadly classified into the three categories:

- 123     ▪ **Unit Testing:** Covers the testing of individual modules and their functionalities.
- 124     ▪ **Integration Testing:** Covers the overall flow of the framework. Different PDEs are solved  
125     with different parameters to check if the accuracy after training is within the acceptable  
126     limits. This ensures that the collection of modules work together as expected.

- 127     ▪ **Compatibility Testing:** The framework is tested with different versions of Python such  
128     as 3.8, 3.9, 3.10, 3.11 and on different versions of OS such as Ubuntu, MacOS and  
129     Windows to ensure the compatibility of the framework across different platforms.

## 130 Acknowledgements

131 We thank Shell Research, India, for providing partial funding for this project. We are thankful  
132 to the MHRD Grant No.STARS-1/388 (SPADE) for partial support.

## 133 References

- 134 Abadi, M., Agarwal, A., Barham, P., Brevdo, E., Chen, Z., Citro, C., Corrado, G. S., Davis,  
135 A., Dean, J., Devin, M., Ghemawat, S., Goodfellow, I., Harp, A., Irving, G., Isard, M., Jia,  
136 Y., Jozefowicz, R., Kaiser, L., Kudlur, M., ... Zheng, X. (2015). *TensorFlow: Large-scale*  
137 *machine learning on heterogeneous systems*. <https://doi.org/10.48550/arXiv.1603.04467>
- 138 Anandh, T., Ghose, D., Jain, H., & Ganesan, S. (2024). *FastVPINNs: Tensor-driven accelera-*  
139 *tion of VPINNs for complex geometries*. <https://doi.org/10.48550/arXiv.2404.12063>
- 140 Ganesan, S., & Tobiska, L. (2017). *Finite elements: Theory and algorithms*. Cambridge  
141 University Press. <https://doi.org/10.1017/9781108235013>
- 142 Kharazmi, E. (2023). *hp-VPINNs: High-performance variational physics-informed neural*  
143 *networks*. <https://github.com/ehsankharazmi/hp-VPINNs>
- 144 Kharazmi, E., Zhang, Z., & Karniadakis, G. E. (2019). Variational physics-informed neural  
145 networks for solving partial differential equations. *arXiv Preprint arXiv:1912.00873*. <https://doi.org/10.48550/arXiv.1912.00873>
- 146
- 147 Kharazmi, E., Zhang, Z., & Karniadakis, G. E. (2021). hp-VPINNs: Variational physics-  
148 informed neural networks with domain decomposition. *Computer Methods in Applied*  
149 *Mechanics and Engineering*, 374, 113547. <https://doi.org/10.1016/j.cma.2020.113547>
- 150 Raissi, M., Perdikaris, P., & Karniadakis, G. E. (2019). Physics-informed neural networks:  
151 A deep learning framework for solving forward and inverse problems involving nonlinear  
152 partial differential equations. *Journal of Computational Physics*, 378, 686–707. <https://doi.org/10.1016/j.jcp.2018.10.045>
- 153